

# Machine Learning

CSCI 5622 Fall 2020

Prof. Claire Monteleoni



# Today

- Discriminative learning II
  - Support vector machine (SVM), continued
  - Kernels
  - Kernel Perceptron

with much credit to S. Dasgupta and T. Jaakkola

# Hard-margin SVM (first version we saw)

- Several desirable properties
  - maximizes the margin on the training set ( $\approx$  good generalization)
  - the solution is unique and sparse ( $\approx$  good generalization)
- But...
  - the solution is very sensitive to outliers, and labeling errors, as they may drastically change the resulting max-margin boundary
  - if the training set is not linearly separable, there's no solution!

# Soft-margin SVM

- We relax the optimization problem by adding slack variables
- Now, not all the constraints need to be met
- The solution therefore need not:
  - Classify all training points with a margin
  - Correctly classify all training points
- The margin is still the region within  $\frac{1}{\|\underline{\theta}^*\|}$  of the decision boundary


# Soft-margin SVM

- Relaxed quadratic optimization problem

$$\text{minimize } \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to}$$

$$y_i(\underline{\theta} \cdot \underline{x}_i + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n$$

$$\xi_i \geq 0, \quad i = 1, \dots, n$$




# Soft-margin SVM

- Relaxed quadratic programming problem

$$\underset{\underline{\theta}, \theta_0, \xi}{\text{minimize}} \quad \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to}$$

$$y_i(\underline{\theta} \cdot \underline{x}_i + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n$$

$$\xi_i \geq 0, \quad i = 1, \dots, n$$



Some of the points may now violate the margin constraints (positive slack) or even be misclassified


# Soft-margin SVM

- The solution now has three types of support vectors

$$\text{minimize } \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to}$$

$$y_i(\underline{\theta} \cdot \underline{x}_i + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n$$

$$\xi_i \geq 0, \quad i = 1, \dots, n$$




# Soft-margin SVM

- The solution now has three types of support vectors

$$\text{minimize } \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to}$$

$$y_i(\underline{\theta} \cdot \underline{x}_i + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n$$

$$\xi_i \geq 0, \quad i = 1, \dots, n$$






# Soft-margin SVM

- The solution now has three types of support vectors

$$\text{minimize } \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to}$$

$$y_i(\underline{\theta} \cdot \underline{x}_i + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n$$

$$\xi_i \geq 0, \quad i = 1, \dots, n$$



# Soft-margin SVM

- Relaxed quadratic programming problem

$$\underset{\underline{\theta}, \theta_0, \xi}{\text{minimize}} \quad \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to}$$

$$y_i(\underline{\theta} \cdot \underline{x}_i + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n$$


$$\xi_i \geq 0, \quad i = 1, \dots, n$$

In the solution, we will have either

$$y_i(\underline{\theta}^* \cdot \underline{x}_i + \theta_0^*) = 1 - \xi_i^*, \quad \xi_i^* \geq 0$$

$$y_j(\underline{\theta}^* \cdot \underline{x}_j + \theta_0^*) > 1, \quad \xi_j^* = 0$$

(all the active constraints are SVs)



# Soft-margin SVM

- Relaxed quadratic programming problem

$$\underset{\underline{\theta}, \theta_0, \xi}{\text{minimize}} \quad \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to}$$

$$y_i(\underline{\theta} \cdot \underline{x}_i + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n$$

$$\xi_i \geq 0, \quad i = 1, \dots, n$$


In the solution, we will have either

$$y_i(\underline{\theta}^* \cdot \underline{x}_i + \theta_0^*) = 1 - \xi_i^*, \quad \xi_i^* \geq 0$$

$$y_j(\underline{\theta}^* \cdot \underline{x}_j + \theta_0^*) > 1, \quad \xi_j^* = 0$$


(all the active constraints are SVs)

The solution need not be unique  
in terms of  $\theta_0, \xi$




# Soft-margin SVM: Examples

- $C=100$




# Soft-margin SVM: Examples

- $C=10$




# Soft-margin SVM: Examples

- $C=1$




# Soft-margin SVM: Examples

- $C=0.1$



# Soft-margin SVM: Examples


- $C$  potentially affects the solution even in the separable case
- $C = 1$






# Soft-margin SVM: Examples

- $C$  potentially affects the solution even in the separable case
- $C = 0.1$




# Soft-margin SVM: Examples

- $C$  potentially affects the solution even in the separable case
- $C = 0.01$




# Beyond linear classifiers...

- Many problems are not solved well by a **linear** classifier even if we allow misclassified examples (SVM with slack)
- E.g., data from experiments typically involve “clusters” of different types of examples



# Non-linear classifiers

- Many (low dimensional) problems are not solved well by a linear classifier even with slack
- By mapping examples to feature vectors, and maximizing a linear margin in the **feature space**, we obtain non-linear margin curves in the original space
- By using non-linear feature mappings we get more powerful sets of classifiers



# Non-linear feature mappings

- The easiest way to make the classifier more powerful is to add non-linear coordinates to the feature vectors
- The classifier is still linear in the parameters,  $\underline{\vartheta}$ , not inputs,  $\underline{x}$

$$\underline{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \phi(\underline{x}) = \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \underline{x} + \theta_0)$$

linear classifier

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \phi(\underline{x}) + \theta_0)$$

non-linear classifier

# Non-linear feature mappings

- The easiest way to make the classifier more powerful is to add non-linear coordinates to the feature vectors
- The classifier is still linear in the parameters,  $\underline{\theta}$ , not inputs,  $\underline{x}$

$$\underline{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \phi(\underline{x}) = \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \underline{x} + \theta_0)$$

linear classifier

$$\underline{\theta} \cdot \underline{x} + \theta_0 = 0$$

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \phi(\underline{x}) + \theta_0)$$

non-linear classifier

# Non-linear feature mappings

- The easiest way to make the classifier more powerful is to add non-linear coordinates to the feature vectors
- The classifier is still linear in the parameters,  $\underline{\theta}$ , not inputs,  $\underline{x}$

$$\underline{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \phi(\underline{x}) = \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \underline{x} + \theta_0)$$

linear classifier

$$\underline{\theta} \cdot \underline{x} + \theta_0 = 0$$

$$\theta_1 x_1 + \theta_2 x_2 + \theta_0 = 0$$

linear decision  
boundary

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \phi(\underline{x}) + \theta_0)$$

non-linear classifier

# Non-linear feature mappings

- The easiest way to make the classifier more powerful is to add non-linear coordinates to the feature vectors
- The classifier is still linear in the parameters,  $\underline{\theta}$ , not inputs,  $\underline{x}$

$$\underline{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \phi(\underline{x}) = \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \underline{x} + \theta_0)$$

linear classifier

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \phi(\underline{x}) + \theta_0)$$

non-linear classifier

$$\underline{\theta} \cdot \phi(\underline{x}) + \theta_0 = 0$$



# Non-linear feature mappings

- The easiest way to make the classifier more powerful is to add non-linear coordinates to the feature vectors
- The classifier is still linear in the parameters,  $\underline{\vartheta}$ , not inputs,  $\underline{x}$

$$\underline{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \phi(\underline{x}) = \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \underline{x} + \theta_0)$$

linear classifier

$$f(\underline{x}; \underline{\theta}, \theta_0) = \text{sign}(\underline{\theta} \cdot \phi(\underline{x}) + \theta_0)$$

non-linear classifier


$$\underline{\theta} \cdot \phi(\underline{x}) + \theta_0 = 0$$

$$\theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 \sqrt{2} x_1 x_2 + \theta_5 x_2^2 + \theta_0 = 0$$

non-linear decision boundary


# Non-linear feature mappings

- By expanding the feature coordinates, we still have a linear classifier in the new feature coordinates but a non-linear classifier in the original coordinates




# Non-linear feature mappings

- By expanding the feature coordinates, we still have a linear classifier in the new feature coordinates but a non-linear classifier in the original coordinates




# Non-linear feature mappings

- By expanding the feature coordinates, we still have a linear classifier in the new feature coordinates but a non-linear classifier in the original coordinates




# Non-linear feature mappings

- By expanding the feature coordinates, we still have a linear classifier in the new feature coordinates but a non-linear classifier in the original coordinates



# Non-linear feature mappings

- By expanding the feature coordinates, we still have a linear classifier in the new feature coordinates but a **non-linear classifier in the original coordinates**



# Learning non-linear classifiers


- We can apply the same SVM formulation, just replacing the input examples,  $x_i$ , with (higher dimensional) feature vectors,  $\Phi(x_i)$ .

$$\begin{aligned} \text{minimize} \quad & \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{subject to} \\ & y_i(\underline{\theta} \cdot \underline{\phi(x_i)} + \theta_0) \geq 1 - \xi_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad i = 1, \dots, n \end{aligned}$$


- Note that the cost of solving this quadratic programming problem increases with the dimension of the feature vectors (BUT, we will **avoid** this problem by solving the dual instead – still only  $n$  constraints.)

# Problems to resolve


- By using non-linear feature mappings we get more powerful sets of classifiers
- Computational efficiency?
  - the cost of using higher dimensional feature vectors (seems to) increase with the dimension
- Model selection?
  - how do we choose among different feature mappings?



linear features




2nd order features




3rd order features



# Quadratic kernel



# Quadratic kernel



# Non-linear perceptron, kernels

- Non-linear feature mappings can be dealt with more efficiently through their inner products or “kernels”
- We will begin by turning the perceptron classifier with non-linear features into a “kernel perceptron”
- For simplicity, we drop the offset parameter

$$f(\underline{x}; \underline{\theta}) = \text{sign}(\underline{\theta} \cdot \underline{\phi}(\underline{x}))$$

Initialize:  $\underline{\theta} = 0$

For  $t = 1, 2, \dots$  (applied in a sequence or repeatedly over a fixed training set)

if  $y_t(\underline{\theta} \cdot \underline{\phi}(\underline{x}_t)) \leq 0$  (mistake)

$$\underline{\theta} \leftarrow \underline{\theta} + y_t \underline{\phi}(\underline{x}_t)$$

# On perceptron updates

- Each update adds  $y_t \underline{\phi}(\underline{x}_t)$  to the parameter vector
- Repeated updates on the same example simply result in adding the same term multiple times
- We can therefore write the current perceptron solution as a function of how many times we performed an update on each training example

$$\underline{\theta} = \sum_{i=1}^n \alpha_i y_i \underline{\phi}(\underline{x}_i)$$

$$\alpha_i \in \{0, 1, \dots\}, \quad \sum_{i=1}^n \alpha_i = \# \text{ of mistakes}$$

where  $\alpha_i$  is the number of mistakes made on example  $i$ .

# Kernel perceptron

- By subbing in this “count” representation of  $\underline{\vartheta}$ , we can write the perceptron algorithm entirely in terms of inner products between the feature vectors

$$f(\underline{x}; \underline{\theta}) = \text{sign}(\underline{\theta} \cdot \underline{\phi}(\underline{x})) = \text{sign}\left(\sum_{i=1}^n \alpha_i y_i [\underline{\phi}(\underline{x}_i) \cdot \underline{\phi}(\underline{x})]\right)$$

Initialize:  $\alpha_i = 0, i = 1, \dots, n$

Repeat for  $t = 1, \dots, n$

if  $y_t \left( \sum_{i=1}^n \alpha_i y_i [\underline{\phi}(\underline{x}_i) \cdot \underline{\phi}(\underline{x}_t)] \right) \leq 0$  (mistake)

$$\alpha_t \leftarrow \alpha_t + 1$$

# Kernel perceptron

- By subbing in this “count” representation of  $\underline{\vartheta}$ , we can write the perceptron algorithm entirely in terms of inner products between the feature vectors

$$f(\underline{x}; \underline{\theta}) = \text{sign}(\underline{\theta} \cdot \underline{\phi}(\underline{x})) = \text{sign}\left(\sum_{i=1}^n \alpha_i y_i [\underline{\phi}(\underline{x}_i) \cdot \underline{\phi}(\underline{x})]\right)$$

Initialize:  $\alpha_i = 0, i = 1, \dots, n$

Repeat for  $t = 1, \dots, n$

if  $y_t \left( \sum_{i=1}^n \alpha_i y_i [\underline{\phi}(\underline{x}_i) \cdot \underline{\phi}(\underline{x}_t)] \right) \leq 0$  (mistake)

$$\alpha_t \leftarrow \alpha_t + 1$$

# Why inner products?

- For some feature mappings, the inner products can be computed efficiently, **without expanding the feature vectors!**

$$\phi(\underline{x}) \cdot \phi(\underline{x}') = \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix} \cdot \begin{bmatrix} x'_1 \\ x'_2 \\ x'^2_1 \\ \sqrt{2}x'_1x'_2 \\ x'^2_2 \end{bmatrix}$$

$$= (x_1x'_1) + (x_2x'_2) + (x_1x'_1)^2 + 2(x_1x'_1)(x_2x'_2) + (x_2x'_2)^2$$

$$= (x_1x'_1 + x_2x'_2) + (x_1x'_1 + x_2x'_2)^2$$

$$= (\underline{x} \cdot \underline{x}') + (\underline{x} \cdot \underline{x}')^2$$

# Why inner products?

- Instead of explicitly constructing feature vectors, we can try to evaluate their inner product or “kernel” directly.

$$\begin{aligned}\phi(\underline{x}) \cdot \phi(\underline{x}') &= \begin{bmatrix} ? \\ ? \end{bmatrix} \cdot \begin{bmatrix} ? \\ ? \end{bmatrix} \\ &= (\underline{x} \cdot \underline{x}') + (\underline{x} \cdot \underline{x}')^2\end{aligned}$$

- What is  $\phi(\underline{x})$  such that the above holds?



# Why inner products?

- Instead of explicitly constructing feature vectors, we can try to explicate their inner product or “kernel”

$$\phi(\underline{x}) \cdot \phi(\underline{x}') = \begin{bmatrix} ? \\ \end{bmatrix} \cdot \begin{bmatrix} ? \\ \end{bmatrix}$$

$$= (\underline{x} \cdot \underline{x}') + (\underline{x} \cdot \underline{x}')^2 + (\underline{x} \cdot \underline{x}')^3 + (\underline{x} \cdot \underline{x}')^4$$

- What is  $\phi(\underline{x})$  now? Does it even exist?

# Feature mappings and kernels

- In the kernel perceptron algorithm, the feature vectors appear only as inner products
- Instead of explicitly constructing feature vectors, we can try to evaluate their inner product or kernel
- $K : \mathcal{R}^d \times \mathcal{R}^d \rightarrow \mathcal{R}$  is a kernel function if there exists a feature mapping such that:

$$K(\underline{x}, \underline{x}') = \underline{\phi}(\underline{x}) \cdot \underline{\phi}(\underline{x}')$$

# Feature mappings and kernels

- In the kernel perceptron algorithm, the feature vectors appear only as inner products
- Instead of explicitly constructing feature vectors, we can try to explicate their inner product or kernel
- $K : \mathcal{R}^d \times \mathcal{R}^d \rightarrow \mathcal{R}$  is a kernel function if there exists a feature mapping,  $\Phi(x)$ , such that

$$K(\underline{x}, \underline{x}') = \underline{\phi}(\underline{x}) \cdot \underline{\phi}(\underline{x}')$$

- Examples of **polynomial** kernels

$$K(\underline{x}, \underline{x}') = (\underline{x} \cdot \underline{x}')$$

$$K(\underline{x}, \underline{x}') = (\underline{x} \cdot \underline{x}') + (\underline{x} \cdot \underline{x}')^2$$

$$K(\underline{x}, \underline{x}') = (\underline{x} \cdot \underline{x}') + (\underline{x} \cdot \underline{x}')^2 + (\underline{x} \cdot \underline{x}')^3$$

$$K(\underline{x}, \underline{x}') = (1 + \underline{x} \cdot \underline{x}')^p, \quad p = 1, 2, \dots$$

# Polynomial decision surfaces

To get a decision surface which is an arbitrary polynomial of order  $p$ :



Let  $\Phi(\mathbf{x})$  consist of all terms of order  $\leq p$ , such as  $\mathbf{x}_1\mathbf{x}_2\mathbf{x}_3^{p-3}$ .

$$k(\mathbf{x}, \mathbf{z}) = \Phi(\mathbf{x}) \cdot \Phi(\mathbf{z}) = (1 + \mathbf{x} \cdot \mathbf{z})^p$$

# Kernel Perceptron recap

Learning in the higher-dimensional feature space:

```
w = 0
while some y(w · Φ(x)) ≤ 0:
    w = w + y Φ(x)
```

Everything works as before; final  $\mathbf{w}$  is a weighted sum of various  $\Phi(\mathbf{x})$ .

**Problem:** number of features has now increased dramatically.  
OCR data: from 784 to 307,720!

# The kernel trick

[Aizenman, Braverman, Rozonoer, 1964]

No need to explicitly write out  $\Phi(\mathbf{x})$  !

The only time we ever access it is to compute a dot product  $\mathbf{w} \cdot \Phi(\mathbf{x})$  .

If  $\mathbf{w} = \mathbf{a}_1 \Phi(\mathbf{x}^{(1)}) + \mathbf{a}_2 \Phi(\mathbf{x}^{(2)}) + \mathbf{a}_3 \Phi(\mathbf{x}^{(3)})$  then  $\mathbf{w} \cdot \Phi(\mathbf{x}) =$  (weighted) sum of dot products, each of the form  $\Phi(\mathbf{x}) \cdot \Phi(\mathbf{x}^{(i)})$  .

Can we compute such dot products without writing out the  $\Phi(\mathbf{x})$  ' s?

# The kernel trick

Polynomial kernel,  $p=2$ :

In 2-d:

$$\begin{aligned}\Phi(\mathbf{x}) \cdot \Phi(\mathbf{z}) &= (1, \sqrt{2}\mathbf{x}_1, \sqrt{2}\mathbf{x}_2, \mathbf{x}_1^2, \mathbf{x}_2^2, \sqrt{2}\mathbf{x}_1\mathbf{x}_2) \cdot (1, \sqrt{2}\mathbf{z}_1, \sqrt{2}\mathbf{z}_2, \mathbf{z}_1^2, \mathbf{z}_2^2, \sqrt{2}\mathbf{z}_1\mathbf{z}_2) \\ &= 1 + 2\mathbf{x}_1\mathbf{z}_1 + 2\mathbf{x}_2\mathbf{z}_2 + \mathbf{x}_1^2\mathbf{z}_1^2 + \mathbf{x}_2^2\mathbf{z}_2^2 + 2\mathbf{x}_1\mathbf{x}_2\mathbf{z}_1\mathbf{z}_2 \\ &= (1 + \mathbf{x}_1\mathbf{z}_1 + \mathbf{x}_2\mathbf{z}_2)^2 \\ &= (1 + \mathbf{x} \cdot \mathbf{z})^2\end{aligned}$$

In  $d$  dimensions:

$$\begin{aligned}\Phi(\mathbf{x}) \cdot \Phi(\mathbf{z}) &= (1, \sqrt{2}\mathbf{x}_1, \dots, \sqrt{2}\mathbf{x}_d, \mathbf{x}_1^2, \dots, \mathbf{x}_d^2, \sqrt{2}\mathbf{x}_1\mathbf{x}_2, \sqrt{2}\mathbf{x}_1\mathbf{x}_3, \dots, \sqrt{2}\mathbf{x}_{d-1}\mathbf{x}_d) \cdot \\ &\quad (1, \sqrt{2}\mathbf{z}_1, \dots, \sqrt{2}\mathbf{z}_d, \mathbf{z}_1^2, \dots, \mathbf{z}_d^2, \sqrt{2}\mathbf{z}_1\mathbf{z}_2, \sqrt{2}\mathbf{z}_1\mathbf{z}_3, \dots, \sqrt{2}\mathbf{z}_{d-1}\mathbf{z}_d) \\ &= (1 + \mathbf{x}_1\mathbf{z}_1 + \mathbf{x}_2\mathbf{z}_2 + \dots + \mathbf{x}_d\mathbf{z}_d)^2 \\ &= (1 + \mathbf{x} \cdot \mathbf{z})^2\end{aligned}$$

Computing dot products in the 307,720-dimensional feature space takes time proportional to just 784, the original dimension!

Never need to write out  $\Phi(\mathbf{x})$ .

Need  $\mathbf{w}$  – but since it's a linear combination of (kernelized) data points, just store the coefficients.

# Kernel trick

## Why does it work?

1. The only time we ever use the data is to compute dot products  $\mathbf{w} \cdot \Phi(\mathbf{x})$ .
2. And  $\mathbf{w}$  itself is a linear combination of  $\Phi(\mathbf{x})$ 's. If  $\mathbf{w} = \mathbf{a}_1 \Phi(\mathbf{x}^{(1)}) + \mathbf{a}_{22} \Phi(\mathbf{x}^{(22)}) + \mathbf{a}_{37} \Phi(\mathbf{x}^{(37)})$  store it as  $[(1, \mathbf{a}_1), (22, \mathbf{a}_{22}), (37, \mathbf{a}_{37})]$
3. Dot products  $\Phi(\mathbf{x}) \cdot \Phi(\mathbf{z})$  can be computed very efficiently.



# Valid kernels: composition rules

- We can construct valid kernels from simple components
- For any function  $f : R^d \rightarrow R$ , if  $K_1$  is a kernel, then so is

$$1) \quad K(\underline{x}, \underline{x}') = f(\underline{x})K_1(\underline{x}, \underline{x}')f(\underline{x}')$$

- The set of kernel functions is closed under addition and multiplication: if  $K_1$  and  $K_2$  are kernels, then so are

$$2) \quad K(\underline{x}, \underline{x}') = K_1(\underline{x}, \underline{x}') + K_2(\underline{x}, \underline{x}')$$

$$3) \quad K(\underline{x}, \underline{x}') = K_1(\underline{x}, \underline{x}')K_2(\underline{x}, \underline{x}')$$

- The composition rules are also helpful in verifying that a kernel is valid (i.e., corresponds to an inner product of some feature vectors)

# Radial basis kernel

- The feature “vectors” corresponding to kernels may also be infinite dimensional (i.e., functions)
- This is the case, e.g., for the radial basis kernel

$$K(\underline{x}, \underline{x}') = \exp \left( -\beta \|\underline{x} - \underline{x}'\|^2 \right), \quad \beta > 0$$

- Any distinct set of training points, regardless of their labels, are separable using this kernel function!

# Radial basis function (RBF) kernel

- The feature “vectors” corresponding to kernels may also be infinite dimensional (i.e., functions)
- This is the case, e.g., for the radial basis kernel

$$K(\underline{x}, \underline{x}') = \exp \left( -\beta \|\underline{x} - \underline{x}'\|^2 \right), \quad \beta > 0$$

- Any distinct set of training points, regardless of their labels, are **separable** using this kernel function!
- We can use the composition rules to show that this is indeed a valid kernel

$$\begin{aligned} \exp\{-\beta\|\underline{x} - \underline{x}'\|^2\} &= \exp\{-\beta\underline{x} \cdot \underline{x} + 2\beta\underline{x} \cdot \underline{x}' - \beta\underline{x}' \cdot \underline{x}'\} \\ &= \overbrace{\exp\{-\beta\underline{x} \cdot \underline{x}\}}^{f(\underline{x})} \exp\{2\beta\underline{x} \cdot \underline{x}'\} \overbrace{\exp\{-\beta\underline{x}' \cdot \underline{x}'\}}^{f(\underline{x}')} \\ &= f(\underline{x}) \left( 1 + 2\beta(\underline{x} \cdot \underline{x}') + \dots \right) f(\underline{x}') \end{aligned}$$

# Valid kernels

- A kernel function is valid (is a kernel) if there exists some feature mapping such that

$$K(\underline{x}, \underline{x}') = \underline{\phi}(\underline{x}) \cdot \underline{\phi}(\underline{x}')$$

- We can verify this, e.g., via the composition rules
- Equivalently, a kernel is valid if it is symmetric and **for all training sets**, the Gram matrix:

$$\begin{bmatrix} K(\underline{x}_1, \underline{x}_1) & \cdots & K(\underline{x}_1, \underline{x}_n) \\ \vdots & \ddots & \vdots \\ K(\underline{x}_n, \underline{x}_1) & \cdots & K(\underline{x}_n, \underline{x}_n) \end{bmatrix}$$

is positive semi-definite.

# Kernel functions

As one varies  $\Phi$ , what kinds of similarity measures  $\mathbf{K}$  are possible?

Any  $\mathbf{K}$  which satisfies a technical condition (positive semi-definiteness) will correspond to some embedding  $\Phi(\mathbf{x})$ .

So: don't worry about  $\Phi$  and just pick a similarity measure  $\mathbf{K}$  which suits the data at hand.

Popular choice: *Gaussian kernel* (typical choice for RBF)

$$\mathbf{k}(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2 / s^2)$$

# Kernel perceptron revisited

- We can now apply the kernel perceptron algorithm **without ever expanding the feature vectors!**

$$f(\underline{x}; \alpha) = \text{sign}\left( \sum_{i=1}^n \alpha_i y_i \underline{K}(\underline{x}_i, \underline{x}) \right)$$

Initialize:  $\alpha_i = 0, i = 1, \dots, n$

Repeat for  $t = 1, \dots, n$

if  $y_t \left( \sum_{i=1}^n \alpha_i y_i \underline{K}(\underline{x}_i, \underline{x}_t) \right) \leq 0$  (mistake)

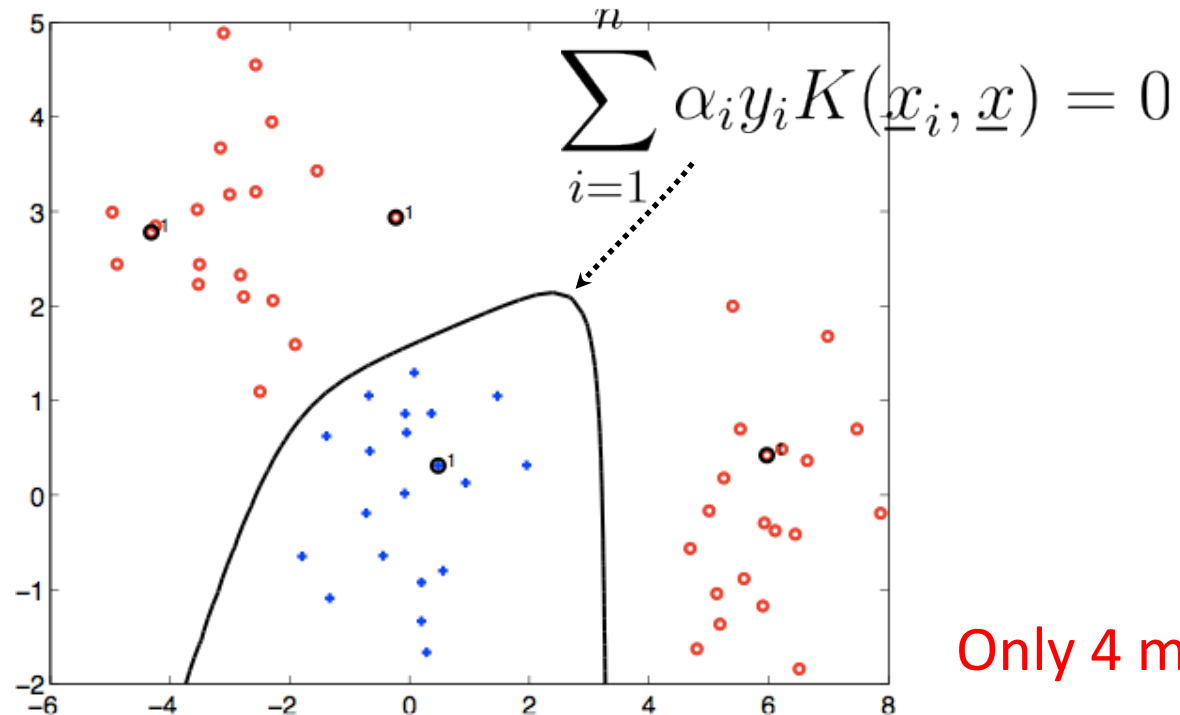
$$\alpha_t \leftarrow \alpha_t + 1$$

# Kernel perceptron: example

- With a radial basis kernel

$$f(\underline{x}; \alpha) = \text{sign}\left(\sum_{i=1}^n \alpha_i y_i K(\underline{x}_i, \underline{x})\right)$$

Decision surface:



Only 4 mistakes!

# Kernel SVM

- Kernel SVM: implicitly find the max-margin linear separator in the feature space, e.g., corresponding to the radial basis kernel

$$f(\underline{x}; \alpha) = \text{sign}\left(\sum_{i=1}^n \alpha_i y_i K(\underline{x}_i, \underline{x}) + \theta_0\right)$$

